

# EMNLP 202

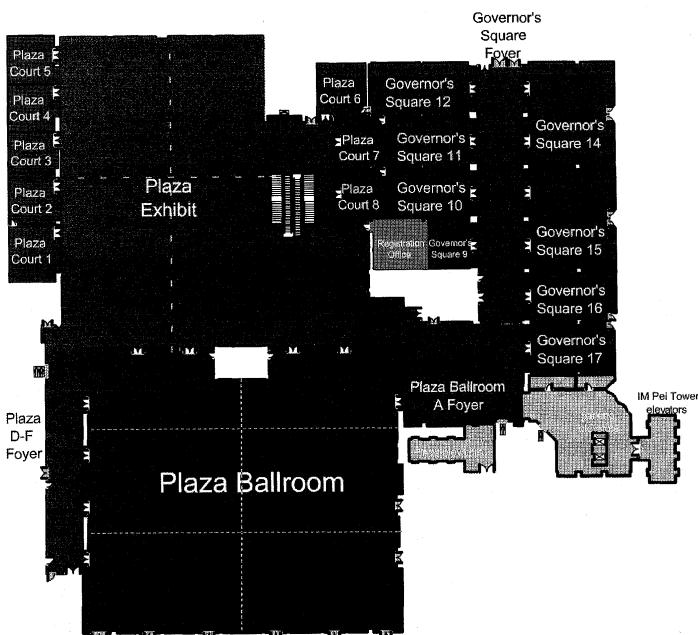


7th - 11th November  
Online & in the Dominican Republic

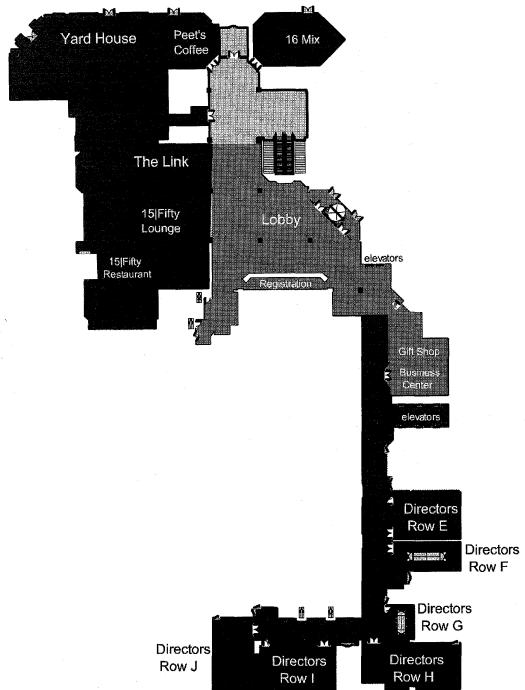
The 2021 Conference on  
Empirical Methods in  
Natural Language Processing

---

**CONFERENCE HANDBOOK**



**CONCOURSE LEVEL  
LOBBY LEVEL**



*Cover by Nathan Cornille  
Handbook assembled by Els Lefever, Pranaydeep Singh and Loic De Langhe  
Printing by Omnipress of Madison, Wisconsin*

---

---

---

## Contents

<b>Table of Contents</b>	<b>i</b>
<b>1 Conference Information</b>	<b>3</b>
Message from the General Chair . . . . .	3
Message from the Program Committee Co-Chairs . . . . .	5
Organizing Committee . . . . .	9
Program Committee . . . . .	10
Meal Info . . . . .	17
<b>2 Tutorials: Wednesday, November 10</b>	<b>19</b>
T1: Crowdsourcing Beyond Annotation: Case Studies in Benchmark Data Collection . . . . .	20
T2: Knowledge-Enriched Natural Language Generation . . . . .	23
<b>3 Tutorials: Thursday, November 11</b>	<b>25</b>
T3: Multi-Domain Multilingual Question Answering . . . . .	26
T4: Robustness and Adversarial Examples in Natural Language Processing . . . . .	27
T5: Syntax in End-to-End Natural Language Processing . . . . .	28
Welcome Reception . . . . .	29
<b>4 Workshops: Thursday–Friday, November 10–11</b>	<b>31</b>
W6: CI+NLP: First Workshop on Causal Inference in NLP . . . . .	33
W10: The Second Workshop on Insights from Negative Results in NLP . . . . .	34
W20: Fourth Workshop on Fact Extraction and VERification (FEVER) . . . . .	35
W21: Third Workshop on New Frontiers in Summarization (NewSum) . . . . .	36
W1: 3rd Workshop on Economics and Natural Language Processing (ECONLP) . . . . .	38
W17: Workshop on Evaluations and Assessments of Neural Conversation Systems . . . . .	40
<b>5 Anti-harassment policy</b>	<b>41</b>
<b>Author Index</b>	<b>43</b>

---





## Big Data

Natural Language Processing • Graph Analytics • Statistical Modeling

Learn more at [goldmansachs.com](http://goldmansachs.com)



© 2015 Goldman Sachs. All rights reserved.

## YAHOO! LABS

### Science-Driven Innovation

Powering products at massive scale with cutting-edge research

Yahoo Labs powers Yahoo's most critical products with innovative science.

- We serve as Yahoo's research arm—its incubator for bold new ideas and laboratory for rigorous experimentation.
- Yahoo Labs applies its scientific findings in powering products for Yahoo's users and enhancing value for its partners and advertisers.
- We create technology that is used by hundreds of millions of customers every day.
- Our forward-looking innovation also helps position Yahoo as an industry and scientific thought leader.

 [labs.yahoo.com](http://labs.yahoo.com)



Join Us!

[SDL.com/Research](http://SDL.com/Research)

Los Angeles, CA

Cambridge, UK

## WHAT WE

**MAKE**

DESCRIBE

INTERPRET

IMPLEMENT

REPORT

DEVELOP

FIGURE OUT

DECODE

FOCUS ON

THINK

POWER

**BREAKS**

MOTIVATES

ENGAGES

INFLUENCES

ADVANCES

DRIVES

SHIFTS

MAXIMIZES

PROMOTES

ENHANCES

ENVIRONMENTS

POSSIBILITIES

OPINIONS

CULTURES

WEALTH

COMMUNITIES

MARKETS

ECONOMIES

OPINIONS

CULTURES

WEALTH

ENVIRONMENTS

POSSIBILITIES

OPINIONS

CULTURES

WEALTH

ENVIRONMENTS

POSSIBILITIES

OPINIONS

**BOUNDARIES**

# Make your mark.

The world's top investors, traders and leaders depend on our information and news. We need your ideas and passion to help us meet some of the biggest challenges around.

[jobs.bloomberg.com](http://jobs.bloomberg.com)

**Bloomberg**



©2015 Bloomberg L.P. All rights reserved. S571691131 0315

## Conference Information

### Message from the General Chair

EMNLP 2021 is one of the first hybrid conferences in the field of natural language processing. It is also for us, the organizing team, uncharted domain. Organizing a hybrid conference has felt like organizing two conferences, a virtual one and an in-person one, which seamlessly must work together and with a kind of multi-task objective make the conference experience synergistic and successful both remotely and in person. With this challenge come opportunities. The hybrid format allows remote participation in a conference that is held onsite in Punta Cana, The Dominican Republic, and allows creating a real conference feeling for those who do not want to travel the many miles from the other side of the world and increase their carbon footprint, and for those who have budget restrictions for traveling. We welcome you all!

As in previous years, the purpose of the General Chair's preface is to express thanks to the amazing team of organizing chairs whose heroic efforts made this hybrid conference possible. The organizing team includes:

- The Programme Chairs – Xuanjing Huang, Lucia Specia and Scott Yih – who did a tremendous job to manage the reviewing process and set up an outstanding scientific program.
- The Senior Area Chairs, Area Chairs and Reviewers whose expertise enabled authors to learn from their reviews and to deliver papers that improved on their original submissions.
- The Demonstration Chairs – Heike Adel and Shuming Shi – who selected outstanding demonstrations to complement the program of the main conference.
- The Workshop Chairs – Minlie Huang and Parisa Kordjamshidi – who made a huge effort for organizing hybrid workshops and satellite conferences.
- The Publication Chairs – Loic Barrault, Greg Durrett and Yansong Feng – who met the challenge of identifying and correcting the myriad ways in which papers could be wrongly formatted, and who assembled the result into the conference proceedings.
- The Handbook Chair – Els Lefever – for the timely delivery of handbook information.
- The Publication Chairs of Findings – Gabriel Stanovsky and Tim Van de Cruys – who made it possible that many interesting papers and their findings can be accessed and cited by the public.

## *Message from the General Chair*

---

- The Tutorial Chairs – Jing Jiang and Ivan Vulic – who selected six excellent tutorials to be presented at the conference.
- The Ethics Chairs – Margot Mieskes and Christopher Potts – who undertook the delicate task of checking papers that had been flagged for potential ethical issues.
- The Website Chair – Miryam de Lhoneux – who ensured that the EMNLP 2021 website promoting this hybrid conference stayed up to date; Mingxiao Li who offered website support; and Nathan Cornille who was responsible for the graphical designs.
- The Virtual Infrastructure Chairs – Quinh Do, Zhaopeng Tu and Dani Yogatama – and the Underline team – Sol Rosenberg, Daniel Luise, Jernej Masnec, Luka Simic, Alexandru Pricop and various support staff.
- The Volunteer Coordinators and Scholarship Chairs – Qi Wu and Diyi Yang – who managed to attract over 200 student and early career volunteers willing to make EMNLP 2021 a success.
- The Publicity Chairs – Raffaella Bernardi and Preethi Jyothi – who have served as both the voice of EMNLP 2021 in communicating with the community and as its ears, reporting on community concerns as soon as they were expressed.
- The Diversity Inclusion Chairs – Laura Alonso Alemany and Toshiaki Nakazawa – who have worked tirelessly to make EMNLP 2021 as welcoming and inclusive as possible for all participants. They have worked with community members to create Birds of a Feather sessions, Affinity Group sessions, student panels and mentoring sessions which contribute to reinforcing the EMNLP community (and sub-groups within this community).

We also want to express special thanks to Priscilla Rasmussen, the ACL Business Manager, first for booking EMNLP 2021 into a beautiful resort in the Dominican Republic, and for the local organization of a hybrid conference. Many thanks, Priscilla!

Finally, we would like to express gratitude to our sponsors, whose generous support has been invaluable in building up EMNLP 2021 to what it is now. These include the Diamond-level sponsors – Apple, Bloomberg Engineering, Facebook AI and Google Research; the Platinum-level sponsors – Amazon Science, Baidu, ByteDance, DeepMind, G Research and Megagon Labs; the Gold-level sponsors – Grammarly and Microsoft; the Silver-level sponsors – duolingo, Naver and Naver Labs Europe; the Bronze-level sponsors – Adobe, Babelscape, human language technology center of excellence and LegalForce; and the Supporter servicenow. We would like to thank the Diversity and Inclusion Champion sponsors – Amazon Science, Deepmind, Google Research and Microsoft; the Diversity and Inclusion Ally sponsor – Morgan Stanley; and the Diversity and Inclusion Contributor sponsors – Adobe and IBM. ACL SIGDAT has also generously contributed to supporting scholarships for attending the conference.

---

*Marie-Francine Moens, KU Leuven, Belgium*  
EMNLP 2021 General Chair

---

## Message from the Program Committee Co-Chairs

---

Welcome to the EMNLP 2021, the first hybrid conference in EMNLP’s history, which is to be held online and in Punta Cana, Dominican Republic.

EMNLP 2021 has received 3,717 full paper submissions, the largest number to date. After excluding papers withdrawn by the authors, and desk rejecting papers which violated the anonymity policy, the multiple submission policy, or the formatting requirements, we were left with 3,600 submissions to be sent out for review. Despite the record-breaking number of submissions, we were able to keep the acceptance rates at a similar level as past years. 841 submissions were accepted to the main conference. Among them, 315 were accepted as oral papers, and 526 were accepted as posters. The decision between oral and poster presentations was not based on the quality/merit of the papers, but on our understanding of what would be the best format for presentation of each individual paper.

We continued providing the acceptance option of “Findings”, following last year’s initiative in the form of a companion publication, for papers that narrowly missed acceptance to the main conference, but were judged to be solid, well-executed research, and worthy of publication. After the review process, 445 papers were invited to be included in the Findings. 26 papers declined the offer, leading to 419 papers to be published in the Findings. Some statistics of the accepted papers are shown below.

	Long	Short	Total
Reviewed	2,540	1,060	3,600
Accepted as Oral	249	66	315
Accepted as Poster	402	124	526
Acceptance Rate (Main Conference)	25.6%	17.9%	23.4%
Accepted to Findings	300	119	419

To meet the reviewer demands of a large conference, we organized the program committee into 22 tracks, including a special “Multidisciplinary and Area Chair Conflict of Interest” track, based on the track information in past conferences. We also introduced a new track called “Efficient methods for NLP” to promote work aiming to reduce the costs of NLP design and experimentation, similar to the “Green NLP” tracks in EACL 2021 and NAACL 2021. In terms of submissions per track, 9 tracks received more than 200 submissions. Particularly popular were the tracks NLP Applications, Machine Learning for NLP, Machine Translation and Information Extraction, which have around 300 submissions each.

We adopted a hierarchical program committee structure similar to that of recent NLP conferences. For each area, we invited 1-4 Senior Area Chair (SACs), who worked with a team of Area Chairs (ACs) they nominated, as well as an army of reviewers that we put together. We used the submission numbers per track from past conferences to estimate the number of SACs and ACs required for each track, leading to 46 SACs and 237 ACs. For reviewer recruitment, we started with the reviewer lists from past conferences and sent out initial invitations asking reviewers to express their track preferences. We then passed the reviewer list to SACs and asked them to select reviewers from these candidate reviewers based on their expertise, and Semantic/Google Scholar profiles. Overall, this resulted in a total of 3,112 reviewers.

Each submission was assigned to three reviewers and one AC. The initial paper assignment was first made using an automatic algorithm to match the abstracts with ACs/reviewers’ past publication records, then adjusted by SACs/PCs. We adapted the review forms from EMNLP 2020, NAACL 2021, and ACL-IJCNLP 2021. Besides the overall recommendation, reviewers were asked to evaluate how reproducible the results in the paper were, and whether there was any ethical concern. Our final decisions were made not just on the review scores, but also took into account the reviews, author responses, discussions among reviewers, meta-reviews and S(AC) recommendations. To ensure the review quality, we provided detailed guidelines about what reviewers should and shouldn’t do in a review.

---

We also formed an Ethics Committee (EC) dedicated to ethical issues. 203 papers with ethical concerns raised by the technical reviewing committee were sent to the EC. The EC chairs went over the papers to determine whether a full EC review would be required. If so, the paper received one or two ethics reviews from additional reviewers recruited by the EC chairs. For any paper that was recommended to be accepted based on technical reviews and that had been referred to the EC, the EC chairs recommended one of the following to the PC chairs: (a) accept (12 EMNLP, 11 Findings), (b) conditionally accept (the ethical issues must be addressed in the camera-ready version; 17 EMNLP, 20 Findings), and (c) reject due to ethical issues (1 paper). The authors of all conditionally accepted papers (except 1 paper declining the Findings offer) submitted the camera-ready version and a short response that explained how they had made the changes requested by the EC meta-reviews. The EC chairs double-checked these revised submissions and responses, and confirmed that the ethical concerns had been addressed. As a result, all conditionally accepted papers were accepted to the main conference or Findings.

ACL Rolling Review (ARR) is a new initiative of the Association for Computational Linguistics, where the reviewing and acceptance of papers to publication venues is done in a two-step process: (1) centralized rolling review and (2) submission to a publication venue. Working closely with the ARR organizers, we ran a pilot at EMNLP 2021. 17 papers (16 long, 1 short) were submitted via ARR to EMNLP 2021, accounting for 25% of the ARR May submissions. After the decision process involving only PCs and SACs, 6 papers (5 long, 1 short) were accepted to the main conference, among which 2 papers were accepted orally. The other 5 long papers were accepted to the Findings. These papers will be published in the respective proceedings as any other EMNLP/Findings paper.

Based on the nominations from SACs and ACs, we identified 21 candidates for the best papers and outstanding papers award. These papers are assessed by the Best Paper Award Committee. The award winners will be announced at the closing ceremony. EMNLP 2021 will also feature 28 papers accepted by the Transactions of the Association for Computational Linguistics (TACL) and 7 papers from the journal of Computational Linguistics (CL), out of which 29 will be presented as orals and 6 as posters. Another highlight of our program is the three exciting keynote talks, presented by Professor Ido Dagan from Bar-Ilan University, entitled “Where next? Towards multi-text consumption via three inspired research lines”, Professor Steven Bird from Charles Darwin University, entitled “LT4All!!? Rethinking the Agenda”, and Professor Evelina Fedorenko from Massachusetts Institute of Technology, entitled “The language system in the human brain”.

There are many people we would like to thank for their significant contributions. EMNLP 2021 would not be possible without their support:

- Our General Chair, Marie-Francine Moens, who has led the whole organizing team, and helped with many of our decision processes;
- 46 SACs who have helped us comprehensively throughout the entire review process, from recruiting ACs and reviewers, assigning papers, checking review quality, making recommendation on final paper decisions, suggesting presentation formats, to recommending best paper candidates; special thanks to Jesse Dodge, who advocated to set up the “Efficient methods for NLP” track, volunteered to serve as the SAC, and helped us update the Reproducibility Checklist to encourage authors to report the computational budget for the experiments in their paper;
- 237 ACs who checked the initial submissions, led paper discussions, wrote meta reviews, ensured review quality, suggested best paper candidates, and recommended outstanding reviewers;

- 3,112 reviewers, 369 secondary reviewers for reviewing papers and actively participating in paper discussions; special thanks to those who stepped in at the last minute to serve as emergency reviewers;
- 35 Ethics Committee members, chaired by Margot Mieskes and Chris Potts, for their hard work to provide ethical reviews and meta-reviews for all papers with serious ethical issues, and ensure that all the conditionally accepted papers have addressed the ethical issues appropriately in a very tight schedule;
- Best Paper Award Committee: Luke Zettlemoyer (chair), Raffaella Bernardi, Mikel L. Forcada, Pascale Fung, Jianfeng Gao, Min Yen KAN, Heng Ji, Mausam, and Ivan Titov, for selecting best papers and outstanding papers under a tight schedule.
- Our postdoc and student assistants Fernando Alva-Manchego, Zichu Fei, Yiding Tan, Yongxin Zhang and Xingwu Hu, who helped with the initial reviewer assignment, anonymity, multiple submission and format checking;
- Past \*ACL PCs, including Trevor Cohn, Yulan He and Yang Liu (EMNLP 2021), Fei Xia, Wenjie Li, RobertoNavigli (ACL-IJCNLP 2021), and Anna Rumshisky, Luke Zettlemoyer and Dilek Hakkani-Tur (NAACL 2021) for all the useful guidance, tips and suggestions on the organization of NLP conferences;
- ARR Editors-in-chief Pascale Fung, Goran Glava?, Sebastian Riedel, Amanda Stent, and CTO Graham Neubig, for their support in running the first ARR pilot, and providing the code for reviewer COI detection and paper assignment;
- Publication Chairs Loic Barrault, Greg Durrett and Yansong Feng, and Findings Chairs Gabriel Stanovsky and Tim Van de Cruys, for completing the final proceedings within a short period;
- ACL Anthology Director Matt Post, for his help in the production of the conference proceedings;
- TACL editors-in-chief Mark Johnson, Ani Nenkova, and Brian Roark, TACL Editorial Assistant Cindy Robinson, and CL Editor-in-Chief Hwee Tou Ng for coordinating TACL and CL presentations with us;
- Workshop Chair Parisa Kordjamshidi and Minlie Huang, for connecting Findings paper authors with workshop organizers for possible presentations.
- Publicity Chairs Raffaella Bernardi and Preethi Jyothi, Website Chair Miryam de Lhoneux, and Website Support Mingxiao Li, who announced conference news on EMNLP Website and social media, collected feedback from the community, and disseminated EMNLP papers with potential public interests via media;
- Rich Gerber at SoftConf, who set up the EMNLP conference site, and was always quick to respond to our emails and resolve any problems we encountered with the START system;
- Sol Rosenberg, Daniel Luise and the whole Underline team, for creating the virtual site for the conference and helping put the hybrid program in place;
- Virtual Infrastructure Chairs Zhaopeng Tu, Dani Yogatama and Quynh Do, who have made the virtual part of the conference possible;
- Els Lefever for preparing the conference handbook;

- Priscilla Rasmussen and members of the Local Organizing Committee, for various discussions on organizing EMNLP, and making the local arrangements for a hybrid conference;
- SIGDAT board members, Iryna Gurevych, Hang Li, Mona Diab and Chin-Yew Lin, for their guidance regarding various decisions;
- The entire EMNLP organizing committee, who have worked together to make EMNLP a success;
- 19,272 authors for submitting their work to EMNLP 2021.

Our deepest gratitude to all of you. We hope you will enjoy the hybrid conference experience.

EMNLP 2021 Program Co-Chairs

*Xuanjing Huang*, Fudan University

*Lucia Specia*, Imperial College London

*Scott Wen-tau Yih*, Facebook

---

# Organizing Committee

---

**General Chair**

Marie-Francine Moens, KU Leuven

**Program Co-chairs**

Xuanjing Huang, Fudan University  
Lucia Specia, Imperial College London  
Scott Wen-tau Yih, Facebook

**Workshop Chairs**

Parisa Kordjamshidi, Michigan State University  
Minlie Huang, Tsinghua University

**Tutorial Chairs**

Jing Jiang, Singapore Management University  
Ivan Vulic, University of Cambridge

**Demonstration Chairs**

Heike Adel, Bosch Center for Artificial Intelligence  
Shuming Shi, Tencent AI lab

**Publication Chairs**

Loic Barrault, University of Sheffield  
Greg Durrett, UT Austin  
Yansong Feng, Peking University

**Findings Chairs**

Gabriel Stanovsky, Hebrew University of Jerusalem  
Tim Van de Cruys, KU Leuven

**Handbook Chair**

Els Lefever, Ghent University

**Publicity Chairs**

Raffaella Bernardi, University of Trento  
Preethi Jyothi, IIT Bombay

**Website Chair**

Miryam de Lhoneux, University of Copenhagen

**Website Support**

Mingxiao Li, KU Leuven

**Virtual Infrastructure Chairs**

Zhaopeng Tu, Tencent AI lab  
Dani Yogatama, Google DeepMind  
Quynh Do, Amazon Aachen

**Local Chair**

Priscilla Rasmussen, ACL Business Manager

**Diversity, Inclusion and Outreach Chairs**

Laura Alonso Alemany, Universidad Nacional de Córdoba  
Toshiaki Nakazawa, The University of Tokyo

**Ethics Committee Chairs**

Margot Mieskes, Darmstadt University of Applied Sciences  
Christopher Potts, Stanford University

**Student Volunteer Coordinator and Scholarship Chairs**

Qi Wu, University of Adelaide  
Diyi Yang, Georgia Tech

---

# Program Committee

---

## Program Committee Co-chairs

Xuanjing Huang, Fudan University  
Lucia Specia, Imperial College London  
Scott Wen-tau Yih, Facebook

## Area Chairs

### *Computational Social Science and Cultural Analytics*

Jing Jiang, Singapore Management University (senior chair)  
Alan Ritter, Georgia Institute of Technology (senior chair)  
Jacob Eisenstein, Google  
Kristen Johnson, Michigan State University  
David Jurgens, University of Michigan  
Preslav Nakov, Qatar Computing Research Institute, HBKU  
Simone Paolo Ponzetto, University of Mannheim  
Andrew Schwartz, Stony Brook University  
Wei Xu, Georgia Institute of Technology  
Justine Zhang, Cornell University

### *Dialogue and Interactive Systems*

Annemarie Friedrich, Bosch Research (senior chair)  
Kevin Small, Amazon (senior chair)  
Ruihua Song, Renmin University (senior chair)  
Emily Dinan, Facebook AI Research  
Zhicheng Dou, Renmin University of China  
Kallirroi Georgila, University of Southern California  
Alborz Geramifard, Facebook AI  
Casey Kennington, Boise State University  
Yanyan Lan, Tsinghua University  
Pearl Pu, EPFL  
Oleg Rokhlenko, Amazon Research  
Gabriel Skantze, KTH Speech Music and Hearing  
Han Ting, National Institute of Advanced Industrial Science and Technology  
Stefan Ultes, Mercedes-Benz AG  
Thang Vu, University of Stuttgart  
Youzheng Wu, JD AI Research  
Wei Wu, meituan  
Min Yang, Chinese Academy of Sciences  
Sina Zarrieß, University of Bielefeld  
Weinan Zhang, Harbin Institute of Technology

### *Discourse and Pragmatics*

Guodong Zhou, Soochow University (senior chair)  
Fang Kong, Soochow University  
Xiaojie Wang, Beijing University of Posts and Telecommunications  
Rui Xia, Nanjing University of Science and Technology  
Ruifeng Xu, Harbin Institute of Technology, Shenzhen

### *Ethics in NLP*

Natalie Schluter, IT University of Copenhagen (senior chair)  
Sasha Luccioni, University of Montreal  
Marta Ruiz Costa-Jussa, Universitat Politècnica de Catalunya  
Simon Suster, University of Melbourne

---

---

### *Efficient Methods for NLP*

Jesse Dodge, (senior chair)

Gabriel Stanovsky, (senior chair)

Aleksandr Drozd, RIKEN Center for Computational Science

Adhi Kuncoro, University of Oxford and DeepMind

Emma Strubell, Carnegie Mellon University

Andreas Rückl, Amazon

Benoît Sagot, Inria

Jonathan Frankle, MIT

### *Generation*

Nan Duan, Microsoft Research Asia (senior chair)

Mohit Iyyer, University of Massachusetts Amherst (senior chair)

Advaith Siddharthan, The Open University (senior chair)

Anya Belz, ADAPT Research Centre, Dublin City University

Asli Celikyilmaz, Facebook AI Research

Nina Dethlefs, University of Hull

Angela Fan, Facebook AI Research

Albert Gatt, Utrecht University

Yeyun Gong, Microsoft Research Asia

Meng Jiang, University of Notre Dame

Chenghua Lin, Department of Computer Science, University of Sheffield

Nanyun Peng, University of California, Los Angeles

Hannah Rashkin, Google Research

Hiroya Takamura, The National Institute of Advanced Industrial Science and Technology

Rui Yan, Renmin University of China

Jiajun Zhang, Institute of Automation Chinese Academy of Sciences

### *Information Extraction*

Leon Derczynski, IT University of Copenhagen (senior chair)

Fei Liu, University of Central Florida (senior chair)

Kang Liu, Institute of Automation, Chinese Academy of Sciences (senior chair)

Christos Christodoulopoulos, Amazon Research

Manuel Ciosici, University of Southern California

Gerard de Melo, Hasso Plattner Institute, University of Potsdam

Antoine Doucet, University of La Rochelle

Doug Downey, Allen Institute for AI, Northwestern University

Mahmoud El-Haj, Lancaster University

Yansong Feng, Peking University

Xianpei Han, Peking University

Luheng He, Google

Zhiyuan Liu, Tsinghua University

Stephen Mayhew, Duolingo

Tim Miller, Boston Children's Hospital and Harvard Medical School

Thien Huu Nguyen, University of Oregon

Qiang Ning, Amazon

Marius Pasca, Google

Hoifung Poon, Microsoft Research

Roi Reichart, Technion - Israel Institute of Technology

Xiang Ren, University of Southern California

Yangqiu Song, HKUST

Jie Yang, Zhejiang University

Mo Yu, IBM Research

Dian Yu, Tencent AI Lab

---

*Information Retrieval and Text Mining*

Mark Sanderson, RMIT University (senior chair)  
Andrew Yates, Max Planck Institute for Informatics (senior chair)  
Simone Filice, amazon.com  
Ahmed Hassan Awadallah, Microsoft Research  
Evangelos Kanoulas, University of Amsterdam  
Sarvnaz Karimi, amazon.com  
Heri Ramampiaro, Norwegian University of Science and Technology (NTNU)  
Suzan Verberne, LIACS, Leiden University  
Thuy Vu, Amazon  
Jenny Zhang, RMIT University

*Interpretability and Analysis of Models for NLP*

Zack Lipton, Carnegie Mellon University (senior chair)  
Francesca Toni, Imperial College London (senior chair)  
Leila Arras, Fraunhofer Heinrich Hertz Institute  
Jasmijn Bastings, Google  
Grzegorz Chrupala, Tilburg University  
Kevin Clark, Stanford University  
Yonatan Belinkov, Technion  
Sebastian Gehrmann, King's College London  
Diewke Hupkes, Facebook AI Research  
Piyawat Lertrvitayakumjorn, Imperial College London  
Anna Rogers, University of Copenhagen  
Naomi Saphra, New York University  
Sameer Singh, University of California, Irvine  
Byron Wallace, Northeastern University  
Oana Cocarascu, King's College London

*Linguistic Theories, Cognitive Modeling and Psycholinguistics*

Ryan Cotterell, ETH Zürich (senior chair)  
Afra Alishahi, Tilburg University  
Kyle Mahowald, University of California, Santa Barbara  
Aida Nematzadeh, DeepMind  
Adina Williams, Facebook, Inc.

*Machine Learning for NLP*

Isabelle Augenstein, University of Copenhagen (senior chair)  
Angeliki Lazaridou, DeepMind (senior chair)  
Wei Lu, Singapore University of Technology and Design (senior chair)  
Karthik Narasimhan, Princeton University (senior chair)  
Yuki Arase, Osaka University  
Daniel Beck, University of Melbourne  
Iz Beltagy, Allen Institute for AI (AI2)  
Kai-Wei Chang, UCLA  
Georgiana Dinu, Amazon AWS  
Lea Fermann, Melbourne University  
Yoav Goldberg, Bar Ilan University  
Yoon Kim, MIT, IBM  
Carolin Lawrence, NEC Laboratories Europe  
Tao Lei, ASAPP Inc  
Lei Li, UCSB  
André Martins, Unbabel, Instituto de Telecomunicacoes

---

---

Vlad Niculae, University of Amsterdam  
Siva Reddy, McGill University  
Vivek Srikumar, University of Utah  
Karl Stratos, Rutgers University  
Jun Suzuki, Tohoku University / RIKEN Center for AIP  
Swabha Swayamdipta, Allen Institute for Artificial Intelligence  
Yulia Tsvetkov, University of Washington  
Sam Wiseman, Toyota Technological Institute at Chicago

*Machine Translation and Multilinguality*

Alexandra Birch, University of Edinburgh (senior chair)  
Yang Feng, Institute of Computing Technology (senior chair)  
Veselin Stoyanov, Facebook (senior chair)  
Boxing Chen, Alibaba  
Colin Cherry, Google  
Trevor Cohn, University of Melbourne  
Marta Ruiz Costa-jussà, Universitat Politècnica de Catalunya  
Marcello Federico, Amazon AI  
Orhan Firat, Google AI  
Dan Garrette, Google Research  
Zhongjun He, Baidu, Inc.  
Tong Xiao, Northeastern University  
Junhui Li, Soochow University, Suzhou  
Yang Liu, Tsinghua University  
Qun Liu, Huawei Noah's Ark Lab  
Rico Sennrich, University of Zurich  
Taro Watanabe, Nara Institute of Science and Technology  
Deyi Xiong, Tianjin University  
Imed Zitouni, Google

*NLP Applications*

Dina Demner-Fushman, National Library of Medicine (senior chair)  
Shafiq Rayhan Joty, Nanyang Technological University (senior chair)  
Maria Liakata, Queen Mary University of London (senior chair)  
Nikolaos Aletras, University of Sheffield  
Emilia Apostolova, Language.ai  
Steven Bedrick, Oregon Health Science University  
Aoife Cahill, DatarMiner  
Nancy Chen, Institute for Infocomm Research, A\*STAR  
Fenia Christopoulou, Huawei Noah's Ark Lab  
Nadir Durrani, QCRI  
Wei Gao, Singapore Management University  
Travis Goodwin, U.S. National Library of Medicine  
Yulan He, University of Warwick  
Lifu Huang, Virginia Tech  
David Mimno, Cornell University  
Makoto Miwa, Toyota Technological Institute  
Tristan Naumann, Microsoft Research  
Dong Nguyen, Utrecht University  
Diarmuid Ó Sághdha, Apple  
Nazneen Rajani, Salesforce Research  
Kirk Roberts, University of Texas Health Science Center at Houston  
Hassan Sajjad, Qatar Computing Research Institute  
Yi Tay, Google  
Thy Thy Tran, Technische Universität Darmstadt

---

---

Karin Verspoor, RMIT University  
Chrysoula Zerva, University of Lisbon  
Aston Zhang, AWS AI  
Arkaitz Zubiaga, Queen Mary University of London

*Phonology, Morphology and Word Segmentation*

Xipeng Qiu, Queen Mary University of London (senior chair)  
Baobao Chao, Institute of Computational Linguistic, Peking University  
Ryan Cotterell, ETH Zürich  
Hinrich Schütze, University of Munich  
Hai Zhao, Shanghai Jiao Tong University  
Tristan Naumann, Microsoft Research

*Question Answering*

Eunsol Choi, UT Austin (senior chair)  
Matt Gardner, Allen Institute for Artificial Intelligence (senior chair)  
Jonathan Berant, Tel Aviv University and AI2  
Jordan Boyd-Graber, University of Maryland  
Danqi Chen, Princeton University  
Yiming Cui, Harbin Institute of Technology  
Hannaneh Hajishirzi, University of Washington  
Robin Jia, Facebook AI Research  
Tushar Khot, Allen Institute for AI  
Tom Kwiatkowski, Google  
Kenton Lee, Google Research  
Jimmy Lin, University of Waterloo  
Minjoon Seo, KAIST  
Pontus Stenetorp, University College London  
Alon Talmor, Allen Institute for AI, Tel-Aviv University

*Resources and Evaluation*

Yvette Graham, ADAPT, Trinity College Dublin (senior chair)  
Barbara Plank, IT University of Copenhagen  
Ines Rehbein, University of Mannheim  
Maja Popovic, ADAPT, Dublin City University  
Gerasimos Lampouras, Huawei Noah's Ark Lab  
Markus Freitag, Google Research  
Simon Mille, Pompeu Fabra University  
Ajay Nagesh, DiDi Labs  
Gareth Jones, Dublin City University

*Semantics: Lexical, Sentence level, Textual Inference and Other areas*

Tim Baldwin, The University of Melbourne (senior chair)  
Sonal Gupta, Facebook (senior chair)  
James Henderson, Idiap Research Institute (senior chair)  
Marianna Apidianaki, University of Helsinki  
Wai Lam, The Chinese University of Hong Kong  
Jey Han Lau, The University of Melbourne  
Mike Lewis, Facebook AI Research  
Koji Mineshima, Keio University  
Nafise Sadat Moosavi, UKP Lab, Technische Universität Darmstadt  
Naoaki Okazaki, Tokyo Institute of Technology  
Tommaso Pasini, University of Copenhagen  
Panupong Pasupat, Google

---

---

Michael Roth, University of Stuttgart  
Swabha Swayamdipta, University of Washington  
Aline Villavicencio, University of Sheffield, UK  
Ivan Vulic, University of Cambridge  
Diyi Yang, Georgia Institute of Technology  
Yi Zhang, Amazon AI

*Sentiment Analysis, Stylistic Analysis, and Argument Mining*

Veronique Hoste, LT3, Ghent University (senior chair)  
Yue Zhang, Westlake University (senior chair)  
Lidong Bing, Alibaba DAMO Academy  
Cristina Bosco, Dipartimento di Informatica - Università di Torino  
Eric Cambria, Nanyang Technological University  
Orphée De Clercq, LT3, Ghent University  
Ivan Habernal, Technische Universität Darmstadt  
Roman Klinger, University of Stuttgart  
Anh Tuan Luu, NTU  
Soujanya Poria, Singapore University of Technology and Design  
Zhiyang Teng, Westlake University  
Zhongqing Wang, Soochow University  
Zhongyu Wei, School of Data Science, Fudan University  
Meishan Zhang, Tianjin University, China

*Speech, Vision, Robotics, Multimodal Grounding*

Hung-yi Lee, Westlake University (senior chair)  
Pranava Madhyastha, City, University of London (senior chair)  
Yonatan Bisk, Carnegie Mellon University  
Christian Fügen, Facebook AI  
David Harwath, The University of Texas at Austin  
Lisa Ann Hendricks, DeepMind  
Chiori Hori, Mitsubishi Electric Research Laboratories (MERL)  
Douwe Kiela, Facebook  
Florian Metze, Carnegie Mellon University  
Tara Sainath, Google, Inc.  
Radu Soricut, Google LLC  
William Wang, University of California, Santa Barbara

*Summarization*

Xiaojun Wan, Peking University (senior chair)  
Lu Wang, University of Michigan (senior chair)  
Giuseppe Carenini, university of british columbia  
Michael Elhadad, Ben Gurion University  
Pengfei Liu, Carnegie Mellon University  
Shashi Narayan, Google  
Manabu Okumura, Tokyo Institute of Technology  
Jessica Ouyang, University of Texas at Dallas  
Maxime Peyrard, EPFL  
Caiming Xiong, Salesforce  
Rui Zhang, Penn State University

*Syntax: Tagging, Chunking and Parsing*

Wanxiang Che, Harbin Institute of Technology (senior chair)  
Liang Huang, Oregon State University and Baidu Research  
Zhenghua Li, Soochow University

---

---

Weiwei Sun, University of Cambridge  
Kewei Tu, ShanghaiTech University  
Anders Søgaard, University of Copenhagen

*Multidisciplinary and AC COI*

Diana Inkpen, University of Ottawa (senior chair)  
Lluís Márquez, Amazon SSAI (senior chair)  
Cecilia Alm, Rochester Institute of Technology  
Paul Cook, University of New Brunswick  
Zornitsa Kozareva, Facebook AI  
German Rigau, UPV/EHU

---

## **Meal Info**

---

The following meals are provided as part of your registration fee:

- A full buffet breakfast will be provided each day in the Plaza Exhibit (foyer)
- Mid-Morning breaks include coffee and tea in the Plaza Exhibit (foyer)
- Mid-Afternoon breaks include coffee, tea, soda, water, and snacks in the Plaza Exhibit (foyer)
- A full dinner buffer is provided during the poster sessions on Monday and Tuesday evenings in the Plaza Exhibit (foyer)

Lunch is provided for students on Monday, but you are otherwise on your own.



# 2

## Tutorials: Wednesday, November 10

### Overview

---

9:00 – 12:30 **Morning Tutorials**

Crowdsourcing Beyond Annotation: Case Studies in Benchmark Data Collection  
*Alane Suhr, Clara Vania, Nikita Nangia, Maarten Sap, Mark Yatskar, Samuel R. Bowman, Yoav Artzi*

Financial Opinion Mining  
*Hsin-Hsi Chen, Hen-Hsen Huang, Chung-Chi Chen*

Knowledge-Enriched Natural Language Generation  
*Wenhao Yu, Meng Jiang, Zhiting Hu, Qingyun Wang, Heng Ji, Nazneen Rajani*

## Tutorial 1

---

**Alane Suhr, Clara Vania, Nikita Nangia, Maarten Sap,  
Mark Yatskar, Samuel R. Bowman, Yoav Artzi**

**Alane Suhr, Clara Vania, Nikita Nangia, Maarten Sap, Mark Yatskar,  
Samuel R. Bowman, Yoav Artzi**

Wednesday, November 10, 2021, 9:00–12:30pm

Crowdsourcing from non-experts is one of the most common approaches to collecting data and annotations in NLP. It has been applied to a plethora of tasks, including question answering, instruction following, visual reasoning, and commonsense reasoning. Even though it is such a fundamental tool, crowdsourcing use is largely guided by common practices and the personal experience of researchers. Developing a theory of crowdsourcing use for practical language problems remains an open challenge. However, there are various principles and practices that have proven effective in generating high quality and diverse data. The goal of this tutorial is to

---

**Alane Suhr** is a PhD student at Cornell University who’s research focuses on grounded natural language understanding. Alane has designed crowdsourcing tasks for collecting language data to study situated natural language understanding. Alane co-presented a tutorial in ACL 2018.

**Clara Vania** is an applied scientist at Amazon. Her research focuses on crowdsourcing, transfer learning, and multilingual NLU. Recently, she has been working on semi-automatic data collection for natural language inference and crowdsourcing methods for question answering.

**Nikita Nangia** is a PhD student at New York University. Nikita’s work focuses on crowdsourcing methods and data creation for natural language understanding. Her recent work explores using incentive structures to illicit creative examples. Nikita co-organized a tutorial on latent structure models for NLP at ACL 2019.

**Maarten Sap** is a PhD student at the University of Washington. His research focuses on endowing NLP systems with social intelligence and social commonsense, and understanding social inequality and bias in language. His substantial experience with crowdsourcing includes the collecting of the SOCIALIQA commonsense benchmark as well as the creation of knowledge graphs with inferential knowledge (ATOMIC, Social Bias Frames).

**Mark Yatskar** is an assistant professor at the University of Pennsylvania. His research focuses on the intersection of natural language processing and computer vision. Mark’s work has resulted in the creation of datasets such as imSitu, QuAC and WinoBias and recent research has focused on gender bias in visual recognition and coreference resolution.

**Sam Bowman** is an assistant professor at New York University. Sam works on data creation, benchmarking, and model analysis for NLU and computational linguistics. Sam has had a substantial role in several NLU datasets, including SNLI, MNLI, XNLI, CoLA, and BLiMP, and his recent work has focused on experimentally evaluating methods for crowdsourced corpus construction.

**Yoav Artzi** is an associate professor at Cornell University. Yoav’s research focuses on learning expressive models for natural language understanding, most recently in situated interactive scenarios. Yoav led tutorials on semantic parsing in ACL 2013, EMNLP 2014 and AAAI 2015.

---

expose NLP researchers to such data collection crowdsourcing methods and principles through a detailed discussion of a diverse set of case studies.

Wednesday, November 10, 2021, 9:00–12:30pm

In this tutorial, we disassemble a financial opinion into 12 components. This tutorial starts by introducing the components one by one and introduces the related studies from both NLP technical aspects and the real-world applications. Besides, in the FinTech trend, financial service gets much attention from the financial industry. However, few studies discuss the opinion toward financial service. In this tutorial, we will also introduce this kind of opinion and provide a comparison with the opinion of investors and customer's opinions in other industries. Several unexplored research questions will be proposed. The audiences of this tutorial will gain an overview of financial opinion mining and figure out their research directions based on the proposed research agenda.

---

**Chung-Chi Chen** is a postdoctoral researcher at the MOST Joint Research Center for AI Technology and All Vista Healthcare, Taiwan. He got the Ph.D. degree in the Department of Computer Science and Information Engineering at National Taiwan University. He received the M.S. degree in Quantitative Finance from National Tsing Hua University, Taiwan. His research focuses on opinion mining and sentiment analysis in finance. He is the organizer of FinNum shared task series in NTCIR (2018–2022) and the FinNLP workshop series in IJCAI (2019–2021). He is the presenter of the ACL-2020 “Natural Language Processing in Financial Technology Applications” tutorial and the presenter of the EMNLP-2021 “Financial Opinion Mining” tutorial. His work has been published in ACL, WWW, SIGIR, IJCAI, and CIKM, and served as PC members in ACL, AAAI, EMNLP, CIKM, and WSDM. He won the 1st prize in both the Jih Sun FinTech Hackathon (2019) and the Standard Chartered FinTech competition (2018), and the 2nd prize in both the Jih Sun FinTech Hackathon (2018) and the E.SUN FHC FinTech Hackathon (2017).

**Hen-Hsen Huang** is an assistant research fellow at the Institute of Information Science, Academia Sinica, Taiwan. His research interests include natural language processing and information retrieval. His work has been published in ACL, SI-GIR, WWW, IJCAI, CIKM, COLING, and so on. Dr. Huang received the Honorable Mention of Doctoral Dissertation Award of ACLCLP in 2014 and the Honorable Mention of Master Thesis Award of ACLCLP in 2008. He served as the registration chair of TAAI 2017, the publication chair of ROCLING 2020, and as PC members of representative conferences in computational linguistics including ACL, COLING, EMNLP, and NAACL. He was one of organizers of FinNum Task at NTCIR and FinNLP Workshop at IJCAI.

**Hsin-Hsi Chen** is a professor in the Department of Computer Science and Information Engineering, National Taiwan University. He was conference chair of IJCNLP 2013, program co-chair of ACM SIGIR 2010, senior PC member of ACM SIGIR 2006, 2007, 2008 and 2009, area/track chair of AAAI 2020, EMNLP 2018, ACL 2012, ACL-IJCNLP 2009 and ACM CIKM 2008, and PC member of many conferences (IJCAI, SIGIR, WSDM, ACL, COLING, EMNLP, NAACL, EACL, IJCNLP, WWW, and so on). He will be conference chair of ACM SIGIR 2023. He received Google research awards in 2007 and 2012, MOST Outstanding Research Award in 2017, and the AmTRAN Chair Professorship in 2018.

---

## Tutorial 2

---

**Wenhao Yu, Meng Jiang, Zhiting Hu, Qingyun Wang,  
Heng Ji, Nazneen Rajani**

**Wenhao Yu, Meng Jiang, Zhiting Hu, Qingyun Wang, Heng Ji, Nazneen Rajani**

Wednesday, November 10, 2021, 9:00–12:30pm

Knowledge-enriched text generation poses unique challenges in modeling and learning, driving active research in several core directions, ranging from integrated modeling of neural representations and symbolic information in the sequential/hierarchical/graphical structures, learning without direct supervisions due to the cost of structured annotation, efficient optimization and inference with massive and global constraints, to language grounding on multiple modalities, and generative reasoning with implicit commonsense knowledge and background knowledge. In this tutorial we will present a roadmap to line up the state-of-the-art methods to tackle these challenges on this cutting-edge problem. We will dive deep into various technical components: how to represent knowledge, how to feed knowledge into a generation model, how to evaluate generation results, and what are the remaining challenges?

---

**Wenhai Yu** is a Ph.D. student in the Department of Computer Science and Engineering at the University of Notre Dame. His research lies in controllable knowledge-driven natural language processing, particularly in natural language generation. His research has been published in top-ranked NLP and data mining conferences such as ACL, EMNLP, AAAI, WWW, and CIKM. Additional information is available at <https://wyu97.github.io/>.

**Meng Jiang** is an assistant professor in the Department of Computer Science and Engineering at the University of Notre Dame. He received his B.E. and Ph.D. in Computer Science from Tsinghua University and was a postdoctoral research associate at the University of Illinois at Urbana-Champaign. His research interests focus on knowledge graph construction and natural language generation for news summarization and forum post generation. The awards he received include Notre Dame Faculty Award in 2019 and Best Paper Awards at ISDSA and KDD-DLG in 2020. Additional information is available at <http://www.meng-jiang.com/>.

**Zhiteng Hu** is an assistant professor in Halıcıoğlu Data Science Institute at UC San Diego. He received his Ph.D. in Machine Learning from Carnegie Mellon University. His research interest lies in the broad area of natural language processing in particular controllable text generation, machine learning to enable training AI agents from all forms of experiences such as structured knowledge, ML systems and applications. His research was recognized with best demo nomination at ACL 2019 and outstanding paper award at ACL 2016. Additional information is available at <http://www.cs.cmu.edu/~zhitengh/>.

**Qingyun Wang** is a Ph.D. student in the Computer Science Department at the University of Illinois at Urbana-Champaign. His research lies in controllable knowledge-driven natural language generation, with a recent focus on the scientific paper generation. He served as a program committee in generation track for multiple conferences including ICML 2020, ACL 2019-2020, ICLR 2021, etc. He previously entered the finalist of the first Alexa Prize competition. Additional information is available at <https://eaglew.github.io/>.

**Heng Ji** is a professor at Computer Science Department of University of Illinois at Urbana-Champaign, and Amazon Scholar. She has published on Multimedia Multilingual Information Extraction and Knowledge-enriched NLG including technical paper generation, knowledge base description, and knowledge-aware image and video caption generation. The awards she received include “Young Scientist” by World Economic Forum, “AI’s 10 to Watch” Award by IEEE Intelligent Systems, NSF CAREER award, and ACL 2020 Best Demo Award. She has served as the Program Committee Co-Chair of many conferences including NAACL-HLT2018, and she is NAACL secretary 2020-2021. Additional information is available at <https://blender.cs.illinois.edu/hengji.html>.

**Nazneen Rajani** is a senior research scientist at Salesforce Research. She got her PhD in Computer Science from UT Austin in 2018. Several of her work has been published in ACL, EMNLP, NACCL, and IJCAI including work on generating explanations for commonsense and physical reasoning. Nazneen was one of the finalists for the VentureBeat Transform 2020 women in AI Research. Her work has been covered by several media outlets including Quanta Magazine, VentureBeat, SiliconAngle, ZDNet. More information on <https://www.nazneenrajani.com>.

---

# 3

## Tutorials: Thursday, November 11

### Overview

---

#### 9:00 – 12:30 Morning Tutorials

Multi-Domain Multilingual Question Answering  
*Sebastian Ruder, Avirup Sil*

Robustness and Adversarial Examples in Natural Language Processing  
*Kai-Wei Chang, He He, Robin Jia, Sameer Singh*

Syntax in End-to-End Natural Language Processing  
*Hai Zhao, Rui Wang, Kehai Chen*

## Tutorial 3

---

### Sebastian Ruder, Avirup Sil

Sebastian Ruder, Avirup Sil

Thursday, November 11, 2021, 9:00–12:30pm

Question answering (QA) is one of the most challenging and impactful tasks in natural language processing. Most research in QA and tutorials, however, has focused on the open-domain or monolingual setting while most real-world applications deal with specific domains or languages. In this tutorial, we attempt to bridge this gap. Firstly, we introduce standard benchmarks in multi-domain and multilingual QA. In both scenarios, we discuss state-of-the-art approaches that achieve impressive performance by either zero-shot learning or out-of-the-box training on open (and closed)-domain QA systems. Finally, we will present open research problems that this new research agenda poses such as multi-task learning, cross-lingual transfer learning, domain adaptation and training large scale pre-trained multilingual language models.

---

**Sebastian Ruder** is a research scientist at DeepMind where he works on transfer learning and multilingual natural language processing. He has been area chair in machine learning and multilinguality for major NLP conferences including ACL and EMNLP and has published papers on multilingual question answering (Artetxe et al., 2020; Hu et al., 2020). He was the Co-Program Chair for EurNLP 2019 and has co-organized the 4th Workshop on Representation Learning for NLP at ACL 2019. He has taught tutorials on “Transfer learning in natural language processing” and “Unsupervised Cross-lingual Representation Learning” at NAACL 2019 and ACL 2019 respectively. He has also co-organized and taught at the NLP Session at the Deep Learning Indaba 2018 and 2019.

**Avirup Sil** is a Research Scientist and the Team Lead for Question Answering in the Multilingual NLP group at IBM Research AI. His team (comprising of research scientists and engineers) works on research on industry scale NLP and Deep Learning algorithms. His team’s system called ‘GAAMA’ has obtained the top scores in public benchmark datasets (Kwiatkowski et al., 2019) and has published several papers on question answering (Chakravarti et al., 2019; Castelli et al., 2020; Glass et al., 2020). He is also the Chair of the NLP professional community of IBM. Avi is a Senior Program Committee Member and the Area Chair in Question Answering for major NLP conferences e.g. ACL, EMNLP, NAACL and has published several papers on Question Answering. He has taught a tutorial at ACL 2018 on “Entity Discovery and Linking”. He has also organized the workshop on the “Relevance of Linguistic Structure in Neural NLP” at ACL 2018. He is also the track coordinator for the Entity Discovery and Linking track at the Text Analysis Conference.

---

## Tutorial 4

---

### Kai-Wei Chang, He He, Robin Jia, Sameer Singh

**Kai-Wei Chang, He He, Robin Jia, Sameer Singh**

Thursday, November 11, 2021, 9:00–12:30pm

Recent studies show that many NLP systems are sensitive and vulnerable to a small perturbation of inputs and do not generalize well across different datasets. This lack of robustness derails the use of NLP systems in real-world applications. This tutorial aims at bringing awareness of practical concerns about NLP robustness. It targets NLP researchers and practitioners who are interested in building reliable NLP systems. In particular, we will review recent studies on analyzing the weakness of NLP systems when facing adversarial inputs and data with a distribution shift. We will provide the audience with a holistic view of 1) how to use adversarial examples to examine the weakness of NLP models and facilitate debugging; 2) how to enhance the robustness of existing NLP models and defense against adversarial inputs; and 3) how the consideration of robustness affects the real-world NLP applications used in our daily lives. We will conclude the tutorial by outlining future research directions in this area.

---

**Kai-Wei Chang** is an assistant professor in the Department of Computer Science at the University of California Los Angeles. His research interests include designing robust, fair, and accountable machine learning methods for building reliable NLP systems (e.g., Alzantot et al., 2018; Shi et al., 2019). His awards include the EMNLP Best Long Paper Award (2017), the KDD Best Paper Award (2010), and the Sloan Research Fellowship (2021). Kai-Wei has given tutorials at NAACL 15, AAAI 16, FAccT18, EMNLP 19, AAAI 20, MLSS 21 on different research topics. Additional information is available at <http://kwchang.net>.

**He He** is an assistant professor in the Department of Computer Science and the Center for Data Science at the New York University. Her research interests include reliable natural language generation and robust learning algorithms that avoid spurious correlations in the data (e.g., He et al., 2019; Tu et al., 2020). She has given tutorials at NAACL 15 and EMNLP 19. Additional information is available at <http://hhexy.github.io>.

**Robin Jia** is currently a visiting researcher at Facebook AI Research, and will be an assistant professor in the Department of Computer Science at the University of Southern California starting in the Autumn of 2021. His research focuses on making natural language processing models robust to unexpected test-time distribution shifts (e.g., Jia and Liang, 2017; Jia et al., 2019). Robin’s work has received an Outstanding Paper Award at EMNLP 2017 and a Best Short Paper Award at ACL 2018. Additional information is available at <https://robinjia.github.io>.

**Sameer Singh** is an Assistant Professor of Computer Science at the University of California, Irvine. He is working on large-scale and interpretable machine learning models for NLP (e.g., Wallace et al., 2019a; Pezeshkpour et al., 2019). His work has received paper awards at ACL 2020, AKBC 2020, EMNLP 2019, ACL 2018, and KDD 2016. Sameer presented the Deep Adversarial Learning Tutorial (Wang et al., 2019) at NAACL 2019 and the Mining Knowledge Graphs from Text Tutorial at WSDM 2018 and AAAI 2017, along with tutorials on Interpretability and Explanations in upcoming NeurIPS 2020 and EMNLP 2020. Sameer has also received teaching awards at UCI. Website: <http://sameersingh.org/>.

---

## Tutorial 5

---

### Hai Zhao, Rui Wang, Kehai Chen

Hai Zhao, Rui Wang, Kehai Chen

Thursday, November 11, 2021, 9:00–12:30pm

This tutorial surveys the latest technical progress of syntactic parsing and the role of syntax in end-to-end natural language processing (NLP) tasks, in which semantic role labeling (SRL) and machine translation (MT) are the representative NLP tasks that have always been beneficial from informative syntactic clues since a long time ago, though the advance from end-to-end deep learning models shows new results. In this tutorial, we will first introduce the background and the latest progress of syntactic parsing and SRL/NMT. Then, we will summarize the key evidence about the syntactic impacts over these two concerning tasks, and explore the behind reasons from both computational and linguistic background.

---

**Hai Zhao** is a professor at the Department of Computer Science and Engineering, Shanghai Jiao Tong University, China. His research interest is natural language processing. He has published more than 120 papers in ACL, EMNLP, COLING, ICLR, AAAI, IJCAI, and IEEE TKDE/TASLP. He won the first place in several NLP shared tasks, such as CoNLL and SIGHAN Bakeoff and top ranking in remarkable machine reading comprehension task leaderboards such as SQuAD2.0 and RACE. He has taught the course “natural language processing” in SJTU for more than 10 years. He is ACL-2017 area chair on parsing, and ACL- 2018/2019 (senior) area chairs on morphology and word segmentation.

**Rui Wang** is a tenured researcher at the Advanced Translation Technology Laboratory, National Institute of Information and Communications Technology (NICT), Japan. His research focuses on machine translation (MT), a classic task in NLP. His recent interests are traditional linguistic based and cutting-edge machine learning based approaches for MT. He (as the first or the corresponding authors) has published more than 30 MT papers in top-tier NLP/ML/AI conferences and journals, such as ACL, EMNLP, ICLR, AAAI, IJCAI, IEEE/ACM transactions, etc. He has also won several first places in top-tier MT shared tasks, such as WMT- 2018, WMT-2019, WMT-2020, etc. He has given several tutorial and invited talks in conferences, such as CWMT, CCL, etc. He served as the area chairs of ICLR-2021 and NAACL- 2021.

**Kehai Chen** is a postdoctoral researcher at the Advanced Translation Technology Laboratory, National Institute of Information and Communications Technology (NICT), Japan. His research focuses on linguistic-motivated machine translation (MT), a classic NLP task in AI. He has published more than 20 MT and NLP papers in top-tier NLP/ML/AI conferences and journals, such as ACL, ICLR, AAAI, EMNLP, IEEE/ACM Transactions on Audio, Speech, and Language Processing, ACM Transactions on Asian and Low-Resource Language Information Processing, etc. He served as a senior program committee of AAAI-2021.

---

## Welcome Reception

---

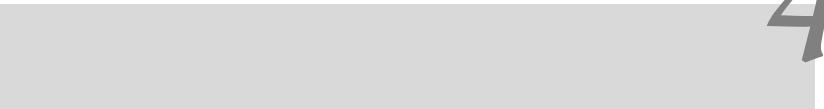
Thursday, November 11, 2021, 6:00pm – 9:00pm

Sheraton Denver Downtown Hotel (conference venue)  
Foyer

Catch up with your colleagues at the **Welcome Reception!** It will be held immediately following the Tutorials on Thursday, November 11 at 6:00pm in the Plaza Exhibit (foyer) of the Sheraton Denver Downtown Hotel (the conference venue). Refreshments and a light dinner will be provided, and a cash bar will be available.



---



4

## Workshops

Wednesday–Thursday

Sixth Conference on Machine Translation	p.??
The 25th Conference on Computational Natural Language Learning (CoNLL)	p.??
8th International Workshop on Argument Mining	p.??
2nd Workshop on Computational Approaches to Discourse (CODI)	p.??

Wednesday

3rd Workshop on NLP for Conversational AI (NLP4ConvAI)	p.??
Novel Ideas in Learning-to-Learn through Interaction	p.??
2nd Workshop on Evaluation and Comparison of NLP Systems	p.??
CI+NLP: First Workshop on Causal Inference in NLP	p.33
The Second Workshop on Insights from Negative Results in NLP	p.34
The 3rd Workshop on Machine Reading for Question Answering	p.??
The Natural Legal Language Processing Workshop 2021 (NLLP)	p.??
SustaiNLP 2021: The Second Workshop on Simple and Efficient Natural Language Processing	p.??
Fourth Workshop on Fact Extraction and VERification (FEVER)	p.35
Third Workshop on New Frontiers in Summarization	p.36

Thursday

Thursday	3rd Workshop on Economics and Natural Language Processing (ECONLP)	p.38
LAW-DMR: The 15th Linguistic Annotation	3rd De-sign-ing Mean-ing Rep-re-sen-ta-tions Joint Work-shop	

p.??

7th Workshop on Noisy User-generated Text (W-NUT 2021) p.??  
The 5th Joint SIGHUM Workshop on Computational p.??

Linguistics for Cultural Heritage, Social Sciences, Humanities, and Literature (LaTeCH-CLfL)

The 1st Workshop on Multi-lines

The 1st Workshop on Multi-lingual Representation Learning (MRL) p...  
The Fifth Widening NLP Workshop (WNLP 2021) p. 23

## Workshop 6: CI+NLP: First Workshop on Causal Inference in NLP

---

Organizers: *Jacob Eisenstein, Amir Feder, Justin Grimmer, Katherine Keith, Emaad Manzoor, Reid Pryzant, Roi Reichart, Molly Roberts, Uri Shalit, Dhanya Sridhar, Brandon Stewart, Victor Veitch, Zach Wood-Doughty, and Diyi Yang*

Venue: Director's Row J

9:00–9:10 **Opening Remarks**

9:10–10:00 **Keynote 1**

10:00–10:30 **Coffee Break**

10:30–12:00 **Short Talks Session 1**

10:50–11:00 Causal Augmentation for Causal Sentence Classification

*Fiona Anting Tan, Devamanyu Hazarika, See-Kiong Ng, Soujanya Poria, and Roger Zimmermann*

11:00–11:10 Text as Causal Mediators: Research Design for Causal Estimates of Differential Treatment of Social Groups via Language Aspects

*Katherine Keith, Douglas Rice, and Brendan O'Connor*

11:10–11:20 Enhancing Model Robustness and Fairness with Causality: A Regularization Approach

*Zhao Wang, Kai Shu, and Aron Culotta*

11:20–11:30 What Makes a Scientific Paper be Accepted for Publication?

*Panagiotis Fytas, Georgios Rizos, and Lucia Specia*

12:00–1:00 **Lunch Break**

1:00–2:30 **Short Talks Session 2**

1:10–1:20 Sensitivity Analysis for Causal Mediation through Text: an Application to Political Polarization

*Graham Tierney and Alexander Volfovsky*

1:20–1:30 A Survey of Online Hate Speech through the Causal Lens

*Antigoni Founta and Lucia Specia*

1:30–1:40 Identifying Causal Influences on Publication Trends and Behavior: A Case Study of the Computational Linguistics Community

*Maria Glenski and Svitlana Volkova*

1:40–1:50 It's quality and quantity: the effect of the amount of comments on online suicidal posts

*Daniel Low, Kelly Zuromski, Daniel Kessler, Satrajit S. Ghosh, Matthew K. Nock, and Walter Dempsey*

2:30–2:45 **Mini Break**

2:45–3:35 **Invited Talk 2**

3:35–4:25 **Invited Talk 3**

4:25–4:45 **Coffee Break**

4:45–5:30 **Panel Discussion**

5:30–6:30 **Poster Session**

---

## Workshop 10: The Second Workshop on Insights from Negative Results in NLP

---

Organizers: *João Sedoc, Anna Rogers, Anna Rumshisky, and Shabnam Tafreshi*

Venue: Governors Square 11

- Corrected CBOW Performs as well as Skip-gram  
*Ozan İrsøy, Adrian Benton, and Karl Stratos*
- Does Commonsense help in detecting Sarcasm?  
*Somnath Basu Roy Chowdhury and Snigdha Chaturvedi*
- BERT Cannot Align Characters  
*Antonis Maronikolakis, Philipp Dufter, and Hinrich Schütze*
- Two Heads are Better than One? Verification of Ensemble Effect in Neural Machine Translation  
*Sungjin Park, Seolhua Lee, Taesun Whang, and Heuiseok Lim*
- Finetuning Pretrained Transformers into Variational Autoencoders  
*Seongmin Park and Jihwa Lee*
- Are BERTS Sensitive to Native Interference in L2 Production?  
*Zixin Tang, Prasenjit Mitra, and David Reitter*
- Zero-Shot Cross-Lingual Transfer is a Hard Baseline to Beat in German Fine-Grained Entity Typing  
*Sabine Weber and Mark Steedman*
- Comparing Euclidean and Hyperbolic Embeddings on the WordNet Nouns Hypernymy Graph  
*Sameer Bansal and Adrian Benton*
- When does Further Pre-training MLM Help? An Empirical Study on Task-Oriented Dialog Pre-training  
*Qi Zhu, Yuxian Gu, Lingxiao Luo, Bing Li, Cheng LI, Wei Peng, Minlie Huang, and Xiaoyan Zhu*
- Recurrent Attention for the Transformer  
*Jan Rosendahl, Christian Herold, Frithjof Petrick, and Hermann Ney*
- On the Difficulty of Segmenting Words with Attention  
*Ramon Sanabria, Hao Tang, and Sharon Goldwater*
- The Highs and Lows of Simple Lexical Domain Adaptation Approaches for Neural Machine Translation  
*Nikolay Bogoychev and Pinzhen Chen*
- Backtranslation in Neural Morphological Inflection  
*Ling Liu and Mans Hulden*
- Learning Data Augmentation Schedules for Natural Language Processing  
*Daphné Chopard, Matthias S. Treder, and Irena Spasić*
- An Investigation into the Contribution of Locally Aggregated Descriptors to Figurative Language Identification  
*Sina Mahdipour Saravani, Ritwik Banerjee, and Indrakshi Ray*
- Blindness to Modality Helps Entailment Graph Mining  
*Liane Guillou, Sander Bijl de Vroe, Mark Johnson, and Mark Steedman*
- Investigating the Effect of Natural Language Explanations on Out-of-Distribution Generalization in Few-shot NLI  
*Yangqiaoyu Zhou and Chenhao Tan*
- Generalization in NLI: Ways (Not) To Go Beyond Simple Heuristics  
*Prajwal Bhargava, Aleksandr Drozd, and Anna Rogers*

- Challenging the Semi-Supervised VAE Framework for Text Classification  
*Ghazi Felhi, Joseph Le Roux, and Djamé Seddah*
- Active Learning for Argument Strength Estimation  
*Natalia Kees, Michael Fromm, Evgeniy Faerman, and Thomas Seidl*

## Workshop 20: Fourth Workshop on Fact Extraction and VERification (FEVER)

---

Organizers: *Rami Aly, Oana Cocarascu, Christos Christodoulopoulos, James Thorne, Andreas Vlachos, and Arpit Mittal*

Venue: Governors Square 11

**Wednesday, November 10, 2021**

**0900–0905 Opening**

**0905–1020 Keynote Talks**

**1020–1100 Research Talks**

10:20–10:35 Evidence Selection as a Token-Level Prediction Task

*Dominik Stammbach*

10:35–10:50 Graph Reasoning with Context-Aware Linearization for Interpretable Fact Extraction and Verification

*Neema Kotonya, Thomas Spooner, Daniele Magazzeni, and Francesca Toni*

11:00–11:15 **Break**

**1115–1230 Keynote Talks**

12:30–1:30 **Lunch**

**1330–1445 Keynote and Shared Task**

2:10–2:20 FaBULOUS: Fact-checking Based on Understanding of Language Over Unstructured and Structured information

*Mostafa BOUZIANE, Hugo Perrin, Amine Sadeq, Thanh Nguyen, Aurélien Cluzeau, and Julien Mardas*

2:20–2:30 Team Papelo at FEVEROUS: Multi-hop Evidence Pursuit

*Christopher Malon*

**1445–1615 Virtual Poster Session**

- Modeling Entity Knowledge for Fact Verification  
*Yang Liu, Chenguang Zhu, and Michael Zeng*

- Verdict Inference with Claims and Retrieved Elements Using RoBERTa  
*In-Zu Gi, Ting-Yu Fang, and Richard Tzong-Han Tsai*

- Stance Detection in German News Articles  
*Laura Mascarell, Tatyana Ruzsics, Christian Schneebeli, Philippe Schlattner, Luca Campanella, Severin Klingler, and Cristina Kadar*

- FANG-COVID: A New Large-Scale Benchmark Dataset for Fake News Detection in German  
*Justus Mattern, Yu Qiao, Elma Kerz, Daniel Wiechmann, and Markus Strohmaier*

- Combining sentence and table evidence to predict veracity of factual claims using TaPaS and RoBERTa  
*Martin Funkquist*

- Automatic Fact-Checking with Document-level Annotations using BERT and Multiple Instance Learning  
*Aalok Sathe and Joonsuk Park*
- Neural Re-rankers for Evidence Retrieval in the FEVEROUS Task  
*Mohammed Saeed, Giulio Alfarano, Khai Nguyen, Duc Pham, Raphael Troncy, and Paolo Papotti*
- A Fact Checking and Verification System for FEVEROUS Using a Zero-Shot Learning Approach  
*Orkun Temiz, Özgün Ozan Kılıç, Arif Ozan Kızıldağ, and Tuğba Taşkaya Temizel*

**1615–1730 Keynote Talks**

5:30–5:45 **Closing Remarks**

## Workshop 21: Third Workshop on New Frontiers in Summarization (NewSum)

---

Organizers: *Lu Wang, Giuseppe Carenini, Jackie Chi Kit Cheung, Fei Liu, and Yue Dong*

Venue: Governors Square 11

9:00–10:30 **Morning Session I**

9:00–9:10 **Open remarks (NewSum Organizers)**

- 9:10–10:00 **Keynote I (Sashi Narayan (Google))**  
10:00–10:10 Sentence-level Planning for Especially Abstractive Summarization  
*Andreas Marfurt and James Henderson*  
10:10–10:20 Template-aware Attention Model for Earnings Call Report Generation  
*Yangchen Huang, Prashant K. Dhingra, and Seyed Danial Mohseni Taheri*  
10:20–10:25 Knowledge and Keywords Augmented Abstractive Sentence  
Summarization  
*Shuo Guan*  
10:25–10:30 Rewards with Negative Examples for Reinforced Topic-Focused  
Abstractive Summarization  
*Khalil Mrini, Can Liu, and Markus Dreyer*

10:30–11:00 **Coffee break I**

11:00–12:00 **Morning session II**

- 11:00–11:50 **Keynote II (Sebastian Gehrmann (Google))**  
11:50–12:00 A Novel Wikipedia based Dataset for Monolingual and Cross-Lingual  
Summarization  
*Mehwish Fatima and Michael Strube*

12:00–1:00 **Lunch break**

1:00–2:30 **Afternoon session I**

- 1:00–1:50 **Keynote III (Asli Celikyilmaz (Facebook AI Research))**  
1:50–1:55 Evaluation of Summarization Systems across Gender, Age, and Race  
*Anna Jørgensen and Anders Søgaard*  
1:55–2:00 Evaluation of Abstractive Summarisation Models with Machine  
Translation in Deliberative Processes  
*Miguel Arana-Catania, Rob Procter, Yulan He, and Maria Liakata*  
2:00–2:10 Capturing Speaker Incorrectness: Speaker-Focused Post-Correction for  
Abstractive Dialogue Summarization  
*Dongyub Lee, Jungwoo Lim, Taesun Whang, Seungwoo Cho,  
Mingun Park, and Heuiseok Lim*  
2:10–2:20 Measuring Similarity of Opinion-bearing Sentences  
*Wenyi Tay, Xiuzhen Zhang, Stephen Wan, and Sarvnaz Karimi*  
2:20–2:30 EASE: Extractive-Abstractive Summarization End-to-End using the  
Information Bottleneck Principle  
*Haoran Li, Arash Einolghozai, Srinivasan Iyer, Bhargavi Paranjape,  
Yashar Mehdad, Sonal Gupta, and Marjan Ghazvininejad*

2:30–3:00 **Coffee break**

---

3:00–3:35 **Afternoon session II**

- 3:00–3:10 Context or No Context? A preliminary exploration of human-in-the-loop approach for Incremental Temporal Summarization in meetings  
*Nicole Beckage, Shachi H Kumar, Saurav Sahay, and Ramesh Manuvinakurike*

- 3:10–3:20 Are We Summarizing the Right Way? A Survey of Dialogue Summarization Data Sets  
*Don Tuggen, Margot Mieskes, Jan Deriu, and Mark Cieliebak*

- 3:20–3:30 Modeling Endorsement for Multi-Document Abstractive Summarization  
*Logan Lebanoff, Bingqing Wang, Zhe Feng, and Fei Liu*

- 3:30–3:35 SUBSUME: A Dataset for Subjective Summary Extraction from Wikipedia Documents  
*Nishant Yadav, Matteo Brucato, Anna Fariha, Oscar Youngquist, Julian Killingback, Alexandra Meliou, and Peter Haas*

3:35–4:15 **EMNLP Finding papers - Summarization**

- 4:00–4:05 Retrieval Augmented Code Generation and Summarization  
*Md Rizwan Parvez, Wasi Ahmad, Saikat Chakraborty, Baishakhi Ray, and Kai-Wei Chang*

4:15–4:45 **Coffee break**

4:45–6:00 **Afternoon session III**

- 4:45–4:55 TLDR9+: A Large Scale Resource for Extreme Summarization of Social Media Posts  
*Sajad Sotudeh, Hanieh Deilamsalehy, Franck Dernoncourt, and Nazli Goharian*

- 4:55–5:00 A New Dataset and Efficient Baselines for Document-level Text Simplification in German  
*Annette Rios, Nicolas Spring, Tannon Kew, Marek Kostrzewa, Andreas Säuberli, Mathias Müller, and Sarah Ebling*

5:00–6:00 **Mentoring Program**

## **Workshop 1: 3rd Workshop on Economics and Natural Language Processing (ECONLP)**

---

Organizers: *Udo Hahn, Veronique Hoste, and Amanda Stent*

Venue: Governors Square 17

**Thursday, November 11, 2021**

9:00–9:30 **Opening remarks and status report on economics NLP (Udo Hahn)**

9:30–10:40 **Session 1: datasets for economic NLP**

9:30–10:00 A Fine-Grained Annotated Corpus for Target-Based Opinion Analysis of Economic and Financial Narratives  
*Jiahui Hu and Patrick Paroubek*

10:00–10:20 EDGAR-CORPUS: Billions of Tokens Make The World Go Round  
*Lefteris Loukas, Manos Fergadiotis, Ion Androutsopoulos, and Prodromos Malakasiotis*

10:20–10:40 The Global Banking Standards QA Dataset (GBS-QA)  
*Kyunghwan Sohn, Sunjae Kwon, and Jaesik Choi*

10:40–10:50 **Short break**

10:50–12:00 **Session 2: transformer-based methodologies in economic NLP**

10:50–11:20 Corporate Bankruptcy Prediction with BERT Model  
*Alex Kim and Sangwon Yoon*

11:20–11:40 Is Domain Adaptation Worth Your Investment? Comparing BERT and FinBERT on Financial Tasks  
*Bo Peng, Emmanuele Chersoni, Yu-Yin Hsu, and Chu-Ren Huang*

11:40–12:00 From Stock Prediction to Financial Relevance: Repurposing Attention Weights to Assess News Relevance Without Manual Annotations  
*Luciano Del Corro and Johannes Hoffart*

12:00–12:45 **Keynote: (Gerard Hoberg)**

12:45–2:00 **Lunch Break**

2:00–3:10 **Session 3: applications in economic NLP**

2:00–2:30 Extracting Economic Signals from Central Bank Speeches  
*Maximilian Ahrens and Michael McMahon*

2:30–2:50 Privacy enabled Financial Text Classification using Differential Privacy and Federated Learning  
*Priyam Basu, Tiasa Singha Roy, Rakshit Naidu, and Zumrut Muftuoglu*

2:50–3:10 Using Word Embedding to Reveal Monetary Policy Explanation Changes  
*Akira Matsui, Xiang Ren, and Emilio Ferrara*

3:10–3:40 **Coffee Break**

3:40–5:00 **Session 4: applications in economic NLP**

3:40–4:10 Effective Use of Graph Convolution Network and Contextual Sub-Tree for Commodity News Event Extraction  
*Meisin Lee, Lay-Ki Soon, and Eu-Gene Siew*

4:10–4:40 Cryptocurrency Day Trading and Framing Prediction in Microblog Discourse

*Anna Paula Pawlicka Maule and Kristen Johnson*

4:40–5:00 To What Extent Can English-as-a-Second Language Learners Read Economic News Texts?

*Yo Ehara*

5:00–5:45 **Discussion Round: Planning for an ECONLP Challenge Competition**

5:45–6:00 **Concluding remarks**

## **Workshop 17: Workshop on Evaluations and Assessments of Neural Conversation Systems**

---

Organizers: *Wei Wei, Bo Dai, Tuo Zhao, Lihong Li, Diyi Yang, Yun-Nung Chen, Y-Lan Boureau, Asli Celikyilmaz, Alborz Geramifard, Aman Ahuja, and Haoming Jiang*

Venue: Director's Row H

- Counterfactual Matters: Intrinsic Probing For Dialogue State Tracking  
*Yi Huang, Junlan Feng, Xiaoting Wu, and Xiaoyu Du*
- GCDF1: A Goal- and Context- Driven F-Score for Evaluating User Models  
*Alexandru Coca, Bo-Hsiang Tseng, and Bill Byrne*
- Unsupervised Testing of NLU models with Multiple Views  
*Radhika Arava, Matthew Trager, Boya Yu, and Mohamed Abdelhady*
- A Comprehensive Assessment of Dialog Evaluation Metrics  
*Yi-Ting Yeh, Maxine Eskenazi, and Shikib Mehri*

## Anti-harassment policy

The open exchange of ideas, the freedom of thought and expression, and respectful scientific debate are central to the aims and goals of the ACL. These require a community and an environment that recognizes the inherent worth of every person and group, that fosters dignity, understanding, and mutual respect, and that embraces diversity. For these reasons, ACL is dedicated to providing a harassment-free experience for all the members, as well as participants at our events and in our programs.

Harassment and hostile behavior are unwelcome at any ACL conference, associated event, or in ACL-affiliated on-line discussions. This includes: speech or behavior that intimidates, creates discomfort, or interferes with a person's participation or opportunity for participation in a conference or an event. We aim for ACL-related activities to be an environment where harassment in any form does not happen, including but not limited to: harassment based on race, gender, religion, age, color, appearance, national origin, ancestry, disability, sexual orientation, or gender identity. Harassment includes degrading verbal comments, deliberate intimidation, stalking, harassing photography or recording, inappropriate physical contact, and unwelcome sexual attention. The policy is not intended to inhibit challenging scientific debate, but rather to promote it through ensuring that all are welcome to participate in shared spirit of scientific inquiry. Vexatious complaints and willful misuse of this procedure will render the complainant subject to the same sanctions as a violation of the anti-harassment policy.

It is the responsibility of the community as a whole to promote an inclusive and positive environment for our scholarly activities. In addition, anyone who experiences harassment or hostile behavior may contact any current member of the ACL Executive Committee or contact Priscilla Rasmussen ([acl@aclweb.org](mailto:acl@aclweb.org)), who is usually available at the registration desk during ACL conferences. Members of the executive committee will be instructed to keep any such contact in strict confidence, and those who approach the committee will be consulted before any actions are taken.

### Implementation

This policy should be posted prominently on all ACL conference and workshop webpages, with a notice of a list of people who can be contacted by community members with concerns or complaints, which will be forwarded to the Professional Conduct Committee for investigation.

Approved by ACL Executive Committee, 2016

Revised by ACL Executive Committee, July 15, 2018

The policy is also available from ACL's main page.



NAACL gratefully acknowledges the following sponsors for their support:

**Diamond**

---



**Bloomberg**

Engineering

FACEBOOK AI

**Google Research**

**Platinum**

---

**amazon | science**

 **DeepMind**



**Megagon Labs**

**Gold**

---

 **grammarly**

 **Microsoft**

**Silver**

---

**NAVER**

**NAVER LABS**  
Europe

## Author Index

---

Bronze

---



LegalForce

Supporter

---



Diversity and Inclusion: Champion

---



Diversity and Inclusion: Ally

---

Morgan Stanley

Diversity and Inclusion: Contributor

---

